

FINE TUNING THE ROTATING SKIP LIST

Itamar Talmon, Tal Shoef

Tel Aviv University

ABSTRACT

Skip list [1] is a linked-list-based structure that became an increasingly popular concurrent alternative to search trees due to its logarithmic complexity and local balancing operation. The Rotating skip list [2] is the fastest concurrent skip list to date. In this paper, we investigate, and try to improve upon, the different heuristics used in the Rotating skip list data structure.

1. INTRODUCTION

1.1. Concurrent Skip Lists

Skip lists are probabilistic search structures which provide improved execution time bounds compared with straightforward binary search trees yet are much simpler to implement than any guaranteed- $O(\log n)$ search structure [1]. A skip list comprises multiple levels, each of which is a linked list. Every skip-list node is present at the lowest level, and probabilistically present in each higher level up to some maximum level that is chosen independently and randomly for each node. This maximum is selected using a random number generator with exponential bias: for example, the probability of inserting into level x is often chosen to be 2^{-x} . Due to its logarithmic complexity, simplicity and local balancing operation, skip list became an increasingly popular concurrent alternative to search trees and the default JDK logarithmic concurrent structures (which is an implementation of the lock-free skip list algorithm suggested by Fraser[3]). Another particularly useful property for parallel skip-list design is that a node can be independently inserted at each level in the list. A node is visible as long as it is linked into the lowest level of the list: insertion at higher levels is necessary only to maintain the property that search time is $O(\log n)$. This led to an improved algorithm [4] that avoids contention hotspots by relying on a maintenance thread to raise, lower and clean-up the towers. This technique lets application threads simply insert an element at the bottom or delete an element by marking it as logically deleted. Yet, until recently, typical skip list did not exploit multicore platforms as well as trees. This has changed when the Rotating skip list was introduced. The Rotating skip list [2] combines the rotation of trees and the uncontended nature of skip lists into a multicore-friendly data structure. The core algorithmic novelty is the use of wheels instead of the

usual towers that are linked together to speedup traversals.

1.2. Rotating Skip List Algorithm

In this section, we present an overview of the rotating skip list key-value store implementation. The rotating skip list differs from traditional skip lists in that it is deterministic and uses wheels, its rotations differ from trees rotations in that they execute in constant time to lower the structure. The resulting algorithm is proven to be linearizable and non-blocking (or lock-free). Refer to [2] for more details, figures, pseudo code and correctness proofs.

1.2.1. Key-value store

Key-value stores offer the basis for indexes to speed up access to data sets. They support the following operations:

- $put(k, v)$ inserts the key-value pair $\langle k, v \rangle$ and returns *true* if k is absent, otherwise return \perp
- $delete(k)$ removes k and its associated value and returns *true* if k is present in the store, otherwise return *false*
- $get(k)$ returns the value associated with key k if k is present in the store, otherwise return *false*

1.2.2. Structure Memory

Locality is achieved by using a rotating array sub-structure, called the wheel, detailed later in a separate section. The structure is a skip list, denoted sl , specified with a set of nodes, including two sentinel nodes. The head node is used as an entry point for all accesses to the data structure, it stores a dummy key that is the lowest of all possible keys. A tail node is used to indicate the end of the data structure, its dummy key is strictly larger than any other possible keys. As in other skip lists, node values are ordered in increasing key order from left to right. The global counter ZERO indicates the index of the first level in the wheels, it is set to 0 initially. The node structure contains multiple fields. It first contains a key-value pair denoted by $\langle k, v \rangle$. Two special values $v = \perp$ and $v = node$ indicate that the node is logically deleted and physically deleted, respectively. A node is first logically deleted before being physically removed and two logically deleted

nodes cannot share the same key k . The level represents the level of this nodes wheel, similar to the level of the corresponding tower in a traditional skip list, it indicates that the node keeps track of successors indexed from 0 to level - 1. The succs is the nodes wheel, it stores the successor pointer at each level of the node. next (resp. prev) is a direct pointer to the next (resp. previous) node that contains the smallest key larger than k (resp. the highest key lower than k). Hence the skip list nodes are all linked through a doubly linked list. This doubly linked list allows to backtrack among preceding nodes if the traversal ends up at a deleted node. Finally, the marker is a special mark used only during physical removal.

1.2.3. Traversal

Each update operation (*put*, *delete*) avoids contention hotspots by localizing the modification to the least contended part of the data structure. All adjustments to the upper levels are sequentially executed by a dedicated maintenance thread, described later, hence allowing a deterministic adjustment of the levels. Any traversal, whether it is for updating or simply searching the structure, is executed from the top of the head node traversing wheels from left to right and levels from top to bottom. Each access looks for the position of some key k in the skip list by starting from the top level of the head down to the bottom level. The *get* function starts by traversing the structure from the skip list head, namely *sl.head*, till the bottom level. It records the value of ZERO at the beginning of the traversal into a local zero variable, sets its starting point to the *set.head* before iterating over each level i from the top level of the skip list, which is also the top level of the head *set.head.level* - 1, to zero. Once the *get* has reached the bottom level, *node* is actually set to the node with the largest key $k' < k$. If this node is physically deleted, the traversal backtracks among deleted nodes and invokes *helpremove* to notify the background thread of the nodes to be removed. Note that the traversal can thus reach a node that is not linked by any succs pointer but only one next pointer. Then it updates *next* to the immediate next node (*node.next*). Once the right position at the bottom level indicated by *next.k* > k is reached, the targeted key k is checked. If it is non logically deleted, the associated value *val* is returned, otherwise \perp is returned to indicate that no key k was present in the key-value store. The *put* function is similar in that it first traverses the structure from top to bottom, backtracks from right to left and help remove deleted nodes. The *put* may find that the node with the key it looks for is logically deleted, in which case it simply needs to logically insert it by setting its value to the appropriate one using a CAS. if the node is found as non logically deleted, then put is unsuccessful and returns *false*. Finally, if the put did not find the targeted key k , it creates a new node *node* with key k and value v that is linked to next and inserts it physically using a CAS. The reason why nodes are logically deleted before

being physically removed is to minimize contention. The delete function marks a node as logically deleted by setting its value to \perp . A separate maintenance thread is responsible for traversing the bottom level of the skip list to clean up the deleted nodes as described in later section. The delete executes like the *put* as it also traverses the structure from top to bottom and backtracks to help remove the deleted nodes. It checks whether a key is absent or if its node is logically or physically deleted in which case it returns *false*. Otherwise, this function logically deletes the node using a CAS to mark it. **A heuristic helps deciding whether the delete should help removing. This happens only when the ratio of logically deleted nodes over non-logically deleted nodes (as communicated by the background thread) reaches 3 after what it pays off.** Whether it helps physically removing or not, the delete returns *true* if the node was logically deleted.

1.2.4. Wheels instead of towers

The wheel size is adjusted without having to mutate a pointer, simply by over provisioning a static array and using modulo arithmetic to adjust the mutable levels. The modulo arithmetic guarantees that increasing the index past the end of an array wraps around to the beginning of the array. This allows lowering to be done in constant-time by simply increasing the index of the lowest logical level of all the arrays without incurring contention. A global variable ZERO is used to keep track of the current lowest level, and when a lowering occurs the lowest index level is invalidated by incrementing the ZERO variable. This causes other threads to stop traversing the index levels before they reach the previous lowest level of pointers. If ZERO is increased above the length of the array this will not compromise the program since the arrays are accessed with modulo arithmetic. To illustrate how wheels improve locality of reference, consider testing to see if the current node next in the traversal has a key greater than the search key. If the wheels were represented using distinct objects, then *next.k* would need to be changed to *next.node.k*, reflecting the fact that the key-value information is being stored in a distinct node object from the index next references. This extra layer of indirection can hurt performance of the skip list, especially since this redirects all traversals.

1.2.5. Background thread

The background (or maintenance) thread executes a loop where it sleeps for 50 microseconds (**a heuristic parameter**). The maintenance thread raises wheels of non-deleted nodes by calling the *raisebottomlevel* function. This raise is compensated with a constant-time lowering specific to the algorithm. This periodic adaptation makes it unnecessary to use the traditional pseudo-random generators as the lowering and raising (the action of increasing the height of a tower) become deterministic. The *lowerskiplist* function discards, in constant-time, the entire bottom level of the skip list by

Figure 10 consists of four line plots arranged in a 2x2 grid, showing the performance (TXS/s) of different methods as a function of the number of threads (4, 8, 16, 32, 64). The methods compared are: dynamic-max-level (blue line with circles), max-level-16 (orange line with squares), max-level-20 (green line with triangles), and max-level-12 (red line with diamonds). The top row shows results for a problem size of 1e8, and the bottom row shows results for a problem size of 1e9. The left column is for u_ratio=20, and the right column is for u_ratio=50. In all cases, the performance increases with the number of threads, and the dynamic-max-level method generally achieves the highest performance.

Problem Size	u_ratio	Threads	dynamic-max-level	max-level-16	max-level-20	max-level-12
1e8	20	4	0.5	0.5	0.5	0.5
		8	0.8	0.8	0.8	0.8
		16	1.5	1.5	1.5	1.5
		32	2.8	2.8	2.8	2.8
		64	4.0	4.0	4.0	4.0
		64	4.2	4.2	4.2	4.2
	50	4	0.5	0.5	0.5	0.5
		8	0.8	0.8	0.8	0.8
		16	1.2	1.2	1.2	1.2
		32	2.1	2.1	2.1	2.1
		64	3.0	3.0	3.0	3.0
		64	3.2	3.2	3.2	3.2
1e9	20	4	0.5	0.5	0.5	0.5
		8	0.8	0.8	0.8	0.8
		16	1.5	1.5	1.5	1.5
		32	2.8	2.8	2.8	2.8
		64	4.0	4.0	4.0	4.0
		64	4.2	4.2	4.2	4.2
	50	4	0.5	0.5	0.5	0.5
		8	0.8	0.8	0.8	0.8
		16	1.2	1.2	1.2	1.2
		32	2.1	2.1	2.1	2.1
		64	3.0	3.0	3.0	3.0
		64	3.2	3.2	3.2	3.2

Figure 10 displays four line plots showing the performance (TXS/s) versus the Number of threads (24, 8, 16, 32, 64) for different levels of deletion (0.5, 1, 5, 10) across two different sizes (1e8 and 1e9) and two different ratios (u_ratio=20 and u_ratio=50). The plots show that performance generally increases with the number of threads, and the performance is relatively consistent across different levels of deletion.

Size	Ratio	Threads	0.5	1	5	10
1e8	20	24	0.5	0.5	0.5	0.5
		8	0.8	0.8	0.8	0.8
		16	1.5	1.5	1.5	1.5
		32	2.5	2.5	2.5	2.5
		64	4.0	4.0	4.0	4.0
1e8	50	24	0.5	0.5	0.5	0.5
		8	0.6	0.6	0.6	0.6
		16	1.2	1.2	1.2	1.2
		32	2.0	2.0	2.0	2.0
		64	2.8	2.8	2.8	2.8
1e9	20	24	0.5	0.5	0.5	0.5
		8	0.8	0.8	0.8	0.8
		16	1.5	1.5	1.5	1.5
		32	2.5	2.5	2.5	2.5
		64	4.0	4.0	4.0	4.0
1e9	50	24	0.5	0.5	0.5	0.5
		8	0.6	0.6	0.6	0.6
		16	1.2	1.2	1.2	1.2
		32	2.0	2.0	2.0	2.0
		64	2.8	2.8	2.8	2.8

Figure 10 displays four line plots showing the performance (TXS/s) versus the Number of threads (24, 8, 16, 32, 64) for different configurations. The plots are arranged in a 2x2 grid, comparing the performance of three methods: level-delete-by-ratio-original (blue line with circles), level-delete-by-ratio-2_3 (orange line with squares), and level-delete-by-ratio-3_4 (green line with triangles).

The top row shows results for 1e8 size with u_ratio=20 (left) and u_ratio=50 (right). The bottom row shows results for 1e8 size with u_ratio=20 (left) and u_ratio=50 (right).

In all plots, the performance (TXS/s) increases as the number of threads increases. The 2_3 and 3_4 methods generally show higher performance than the original method, especially at higher thread counts.

As can be seen on figure BLABLALB, ads

6. EXPERIMENTAL SETTINGS

In this section we describe the multicore machines and the benchmarks used in our experiments as well as the 7 data structure algorithms we compare our rotating skip list against. Multicore machines. We used two different multicore machines to validate our results, one with 2 8-way Intel Xeon E5-2450 processors with hyperthreading...

7. FUTURE RESEARCH

BLIA BALIDSA ADBLADS

8. REFERENCES

- [1] William Pugh, "Skip lists: A probabilistic alternative to balanced trees," in *Workshop on Algorithms and Data Structures*, 1989, vol. August, pp. 437–449.
- [2] Alan Fekete Dick, Ian and Vincent Gramoli, "A skip list for multicore," in *Concurrency and Computation: Practice and Experience*, 2017, vol. 29, p. 4.
- [3] Keir Fraser, "Practical lock-freedom," in *University of Cambridge, Computer Laboratory*, 2004, vol. Feb.
- [4] Gramoli Vincent Crain, Tyler and Michel Raynal, "No hot spot non-blocking skip list," in *Distributed Computing Systems (ICDCS)*, 2013, vol. 33, pp. 196–205.

